SUBHASH

LEAD AZURE DATA PLATFORM ARCHITECT | TECHNICAL LEADERSHIP | CLOUD & LAKEHOUSE

subhash1692@gmail.com | github.com/subhash1692 | 15+ years in data and platform engineering

PROFESSIONAL PROFILE

Azure data platform architect and technical leader with 15+ years in data and technology delivery. Designs enterprise lakehouse operating models that connect data ingestion, Databricks engineering, infrastructure provisioning, identity governance, developer self-service, and controlled releases. Combines reference architecture and standards ownership with hands-on implementation, PR review, production recovery, and cost-aware platform operations.

SELECTED SCOPE & DELIVERY FOOTPRINT

- **Platform scale:** Contributed to an enterprise lakehouse spanning 25+ business domains and 467+ Databricks notebooks, with an associated estate of 3,490+ SQL definitions.
- **Engineering enablement:** Established delivery and PR controls across 400+ pipeline artifacts; implemented 8 Bicep infrastructure modules and multi-stage automated pipelines for provisioning and application releases.

CORE CAPABILITIES

Architecture & data Enterprise Azure platforms, Azure Databricks, ADLS Gen2, ADF, PySpark, Spark SQL, Lakeflow/DLT, Delta Lake, Apache Iceberg, Synapse modernization, Fabric/OneLake.

Platform operations Infrastructure as code (Bicep/ARM), Databricks Asset Bundles, Azure DevOps, self-hosted Linux agents, SHIR, operational runbooks, production monitoring, cost visibility.

Governance & trust Unity Catalog, Entra ID, managed identities/service principals, RBAC, Key Vault, PII classification/masking, audit/lineage, permissions-as-code, secure APIs.

Engineering practices Reusable frameworks, technical design/code reviews, CI/CD release gates, data quality and reconciliation, Python/SQL/Bash/PowerShell, testing, GenAI document pipelines.

PROFESSIONAL EXPERIENCE

Sr Consultant, Data Platform | Lingaro India Pvt Ltd Aug 2021 - Present

Platform Architecture & Technical Leadership

- **End-to-end architecture:** Defined platform patterns from enterprise sources through ADF/ADLS ingestion, Databricks transformation, Unity Catalog governance, and business-facing SQL and BI consumption.
- **Engineering standards:** Established reusable delivery templates, pull-request quality checks, branch promotion rules, code review practices, architecture documentation, and operational runbooks for collaborating engineering teams.
- **Reusable data frameworks:** Designed metadata-driven ingestion and transformation capabilities with schema drift, incremental loads, reconciliation, backfills, and explicit data quality checks; enabled consistent onboarding across domains.
- **Self-service platform:** Developed governed workflows for workspace/user onboarding, compute and permission requests, access approval, lifecycle tracking, and resource consumption visibility.

Cloud Infrastructure, Release Engineering & Security

- **Secure deployment architecture:** Established multi-environment CI/CD using Databricks Asset Bundles, Azure DevOps, idempotent SQL-as-code, versioned application artifacts, approval gates, and automated post-release branch synchronization.
- **Self-hosted deployment infrastructure:** Created and maintained Azure DevOps self-hosted Linux agents, including tool bootstrapping and environment configuration for controlled deployments into network-restricted platforms.
- **Integration Runtime ownership:** Maintained Azure Data Factory Self-hosted Integration Runtime (SHIR) on a quarterly maintenance cycle to sustain reliable hybrid integration operations.
- **Repeatable provisioning:** Designed 8 modular Azure Bicep infrastructure templates, implemented a 5-stage release workflow, and built an 8-stage Databricks workspace bootstrap integrating ARM, RBAC, cluster ACLs, and Unity Catalog.
- **Workspace scale:** Delivered parameterized automation supporting batch provisioning of 50+ resource groups, with environment-specific permissions, identity integration, and pre-deployment checks.
- **Governance & access controls:** Automated Unity Catalog object deployment and least-privilege grants across catalogs, schemas, tables, and external locations; integrated service principals, managed identities, Key Vault, and PII-aware lineage monitoring.

Modernization, Reliability & AI Enablement

- **Lakehouse modernization:** Led technical design and implementation of Synapse stored procedure modernization into Spark/Delta processing; preserved business semantics using reconciliation, incremental logic, and safe reruns.
- **Reliability & cost:** Diagnosed ADF dependency failures, credential/configuration drift, and Databricks post-migration performance regressions; tuned execution and improved visibility into compute usage and resource lifecycle.
- **AI-ready analytics:** Developed LLM-driven document classification, summarization, and field extraction into governed data assets; designed resilient endpoints, prompt configuration, and inference-cost-aware reruns.
- **Analytical services:** Delivered provisioning and CI/CD infrastructure for Durable Functions-based forecasting services; supported Fabric/OneLake/Direct Lake analytics and automated Power BI refresh.

Lead Implementation | Buck Consultants Aug 2019 - Aug 2021

- Led reusable ADF ingestion and Synapse/PolyBase integration delivery; produced team implementation standards, documentation, and onboarding patterns.

Sr Analyst | Conduent Business Services Nov 2011 - Aug 2019

- Developed parameterized cloud data ingestion and SQL automation; supported operational reporting, pipeline support, and continuous engineering process improvement.

Sr Associate | MetLife GOSC Mar 2009 - Nov 2011

- Supported insurance applications, SQL Server operations, configuration management, and production support.

CERTIFICATIONS & EDUCATION

Microsoft Certified: Fabric Data Engineer Associate | Azure Data Engineer Associate | Azure Administrator Associate | Azure Data Fundamentals

Bachelor of Commerce, Delhi University